

SANDESH BHARADWAJ

+971-58-275-0191 | sandesh.bharadwaj97@gmail.com |
www.tensorsofthewall.com | github.com/tensorsofthewall |
linkedin.com/in/sandeshbharadwaj97



SKILLS

Machine Learning & Systems: Large-scale LLM/VLM training; Long-context modeling (24K+ tokens); Multi-node multi-GPU training; Distributed inference; Synthetic data generation
Distributed Training: DDP, FSDP, Tensor/Context/Data Parallelism; memory-efficient training; heterogeneous GPU clusters (NVIDIA & AMD)
LLM Training & Alignment: Supervised fine-tuning (SFT), Direct Preference Optimization (DPO), Preference learning; large-scale evaluation pipelines
Frameworks: PyTorch, PyTorch Lightning, HuggingFace Transformers, vLLM, TensorFlow, LangChain, Apache Spark, Apache Airflow, Intel OpenVINO
Languages & Datastores: Python, C/C++, Bash, Java, MySQL, SQLite, ChromaDB, Elasticsearch
Experimentation & Compute: SLURM; Multi-node training & deployment; Reproducible experiments; Cluster debugging; Performance profiling

PUBLICATIONS

- *WellLogBench: A Domain-Expert Curated Benchmark for Evaluating LLM Reasoning on Subsurface Well Log Data.* S. Shikar*, **S. Bharadwaj*** (equal contribution), et al. **ACL 2026 (under review)**
- Unified Local-Cloud Decision-Making via Residual Reinforcement Learning. K. Sengupta, Z. Shangguan, **S. Bharadwaj** et al. **ECCV 2024**
- Person Re-Identification by Analyzing Dynamic Variations in Gait Sequences. **S. Bharadwaj** et al. **ETCCS 2020**

RESEARCH EXPERIENCE

Mohamed bin Zayed University of Artificial Intelligence | Abu Dhabi, UAE Aug 2025 – Present
Research Engineer I [Full-Time]

- Working on large-scale training, post-training, evaluation, and deployment of domain-specific LLMs and VLMs as part of the *Multimodal Energy GPT* project.
- Trained and post-trained foundation LLMs on long-context data using multi-node, multi-GPU setups (up to 16 nodes), leveraging **context, tensor, and data parallelism, DDP, and FSDP** for memory-efficient scaling.
- Executed large-scale supervised fine-tuning and preference optimization (DPO) for 70B-class models, including training on AMD GPU clusters with nD-parallelism for stable multi-node scaling.
- Performed efficient multi-node inference of 100B+ models (Qwen-235B, LLaMA-109B) to generate large-scale synthetic datasets.
- Adapted distributed training and inference pipelines across heterogeneous GPU clusters (NVIDIA and AMD) to ensure consistent performance across vendor-specific environments.

Lossfunk | Remote June 2025 – Present
Researcher (Initialize Program) [Part-Time]

- Awarded compute sponsorship through Lossfunk's Initialize program to independently pursue research on MIDAS, a modular framework for music source separation.
- Designing novel architecture and loss functions to improve separation quality while enabling scalable experimentation with structured priors and frequency-domain transformations

Human-to-Everything (H2X) Lab, Boston University | Boston, MA March 2023 – March 2025
Visiting Researcher / Graduate Research Assistant - Prof. Eshed Ohn-Bar [Full-Time]

- Enhanced vision-language model performance for autonomous driving by 80% through creation of novel dataset and application of fine-tuning, RAG, and RAFT techniques.
- Collaborated with Red Hat to optimize cloud-edge collaboration algorithms, developing a conditional routing module for dynamic offloading that reduced energy consumption.

- Achieved state-of-the-art performance in CARLA-based navigation metrics, outperforming existing solutions by 35% while enhancing energy efficiency.
- Designed a novel crowd navigation environment and introduced the Ecological Navigation Score metric to prioritize energy-efficient navigation in reinforcement learning tasks.
- Developed visual servoing for multi-robot simulations using RGB and depth sensors, integrated human feedback, and evaluated reinforcement learning algorithms for person-following, significantly improving performance.

Teaching Assistant: ENG EC 444 - Smart and Connected Systems [Part-Time]

Jan 2024 – May 2024

- Assisted with an introductory course on cyber-physical and IoT systems covering microcontroller architecture, real-time OS programming, WPANs, IP gateways, cloud computing, reliability, security, and privacy.
- Supported course logistics, conducted classes, graded assignments, managed lab sessions independently and contributed to developing new assignments.

Robert Bosch GmbH | Bangalore, India

Jan 2020 – June 2020

Research Intern - Corporate Research [Full-Time Co-Op]

- Adapted open-set classification to identify rare objects in vehicular datasets by merging classification models with meta-learning and few-shot algorithms, achieving 56% accuracy on the Open Images subset.
- Developed novel algorithm to predict automobile part degradation and recommend preemptive actions.

Center for Development of Advanced Computing | Kolkata, India

May 2019 – Oct 2019

Research Intern [Full-Time]

- Led and implemented modified approach to person re-identification based on dynamic gait.
- Achieved best accuracy of 91.13% on CASIA-B Gait Dataset (Python, Pandas, Numpy, Scikit-learn).
- Presented research paper at ETCCS 2020, indexed in Springer LNEE Series.

INDUSTRY EXPERIENCE

Ottometric Inc. | Waltham, MA

June 2023 – Aug 2023

Software Engineering Intern [Full-Time]

- Developed data summarization algorithms for the Ottoviz platform by creating a custom fork of the *submodlib* library and applying submodular optimization techniques.
- Enabled efficient training of deep learning models by developing custom features for selecting important 1% subset of images, achieving significant dataset reduction.

Ignitarium Technology Solutions | Bangalore, India

Jan 2022 – July 2022

AI Engineer [Full-Time]

- Improved inference performance on Untether AI's accelerator chips by developing high-performance low-level INT8 kernels in C++.
- Implemented generalized padding for transformation kernels, replacing the non-generalized version, and optimized quantization, performance metrics, and execution cycles, achieving 30% performance gain.
- Led team to streamline workflows, create detailed documentation, and reduce codebase size by 40%.

Synopsys India | Bangalore, India

Jan 2021 – Jan 2022

Data Engineer [Full-Time]

- Optimized existing data extraction and processing pipelines with Apache Spark, improving processing time by 25%.
- Developed BERT-based anomaly detection tool to identify, track, and alert users to critical events requiring intervention, achieving average accuracy of 92% in runtime issue detection.
- Redesigned several existing applications for log analysis into a single unified system in production, improving performance by 30% while enhancing user experience, reliability, and accessibility.

Thermo Fisher Scientific India Pvt. Ltd. | Hyderabad, India

May 2018 – July 2018

Summer Intern [Full-Time]

- Developed proof-of-concept for upgrading existing product on ARM hardware using embedded C++, delivering nearly 2x processing power while reducing manufacturing and development costs by 70%.
- Designed user interface with Qt5 and implemented on Linux Yocto, enabling custom Linux-based systems independent of hardware architecture for seamless user interaction.

PROJECTS & OPEN-SOURCE

Exo – github.com/exo-explore/exo

Jan 2025 – Apr 2025

- Started contributing to exo, a distributed systems framework that allows users to create a cluster using everyday devices for running AI models.
- Merged a PR to improve support for AMD GPUs and boost inference throughput on Windows machines.
- Submitted PR to support llama.cpp as primary backend for exo, bringing Windows support on par with macOS / Linux.

BetterSearch - github.com/tensorsofthewall/BetterSearch

July 2024 – Present

- Developed desktop application integrating LLMs, advanced RAG techniques, and native OS functionality for efficient file search and content indexing, with customizable configurations for optimized CPU and GPU performance.
- Integrated osquery to enable text-to-SQL functionality, providing device management capabilities for users.
- Leverages custom fine-tuned LLM for text-to-SQL generation in local deployments, with support for cloud providers such as OpenAI and Google Gemini for scalability.

VidTune - github.com/tensorsofthewall/VidTune

May 2024 – Aug 2024

- Created generative AI web-app that generates tailored music for videos.
- Designed an intuitive interface and advanced customization tools for genre-specific music integration, tempo adjustment, keyword alignment, and dynamic mixing.

Person Following Robot

April 2023 – June 2023

- Developed and trained imitation learning and reinforcement learning algorithms in simulation for effective person tracking and following.
- Created multi-robot Gazebo simulation with leader, follower, and obstacles, achieving 72% accuracy in precise person detection, tracking, and following (Python, PyTorch, OpenCV, ROS, Gazebo).

Evaluating complex ML streaming pipelines

Feb 2023 – April 2023

- Designed benchmarks for complex ML pipelines to evaluate external and embedded serving frameworks.
- External serving models (TorchServe, TF-Serving, OpenVINO Model Server) outperform embedding serving pipelines (TF-Saved and ONNX) by 200%.

Robot Learning and Vision

Jan 2023 – April 2023

- Designed a modular self-driving pipeline in CARLA consisting of Homography Transformation, Lane Detection, Waypoint Prediction, Lateral and Longitudinal Control modules, improving driving score to 53%.
- Designed and trained a custom CNN through imitation learning in CARLA, achieving a driving score of 32.5% and 54% route completion.

Dynamic Gesture Recognition

Oct 2022 – March 2023

- Developed application for dynamic gesture detection using cameras; one of top 25 teams in OpenCV AI Competition 2022 - Spatial AI Track.
- Trained and evaluated vision models (ResNeXT-101, MViT, 3D-MobileNetV2) on Qualcomm Jester dataset, achieving 17% performance improvement with custom Bi-LSTM architecture.

Autonomous Underwater Vehicle - <https://auviiitdm.github.io/>

Sep 2018 – April 2020

- Led a team to the Singapore AUV Challenge finals in 2019 and 2020.
- Contributed to navigation, vision software, and electronics, guiding the team across all aspects of vehicle development.
- Enhanced object detection and tracking by integrating modified Sea-Thru water-removal algorithm into AUV inference pipeline.

EDUCATION

Boston University

Sep 2022 - May 2024

Master of Science in Computer Science

MS Thesis Advisor - Prof. Eshed Ohn-Bar.

Coursework: Image & Video Computing, Streaming and Event-driven Systems,

Robot Learning & Vision for Navigation, Deep Learning,

Designing Computing Systems for Robotics.

Indian Institute of Information Technology

Design and Manufacturing, Kancheepuram, Chennai, India

July 2015 - May 2020

Dual Degree: Bachelor of Technology in Electronics and Communication.

Master of Technology in Signal Processing and Communication Systems Design.

Coursework: Machine Learning, Digital Signal Processing,

Computer Vision, Data Structures and Algorithms.

Positions of: Music Club Secretary.

Responsibility Autonomous Underwater Vehicle Team Lead.

NTSE Scholar

2013 - 2020

National Talent Search Examination, Government of India

- Awarded to top 1% of students nationwide for academic excellence from high school through graduate studies.

REFERENCES

Hisham Cholakkal | Mohamed bin Zayed University of Artificial Intelligence

Assistant Professor

Department of Computer Vision

hisham.cholakkal@mbzuai.ac.ae

Eshed Ohn-Bar | Boston University

Assistant Professor

Department of Electrical and Computer Engineering

ehonbar@bu.edu

Sahana Prabhu | Robert Bosch Engineering and Business Solutions Pvt. Ltd.

Research Scientist and Technical Architect

Sahana.MuraleedharaPrabhu@in.bosch.com

Kunal Chanda | C-DAC, Kolkata

Scientist-E

kunalchanda@cdac.in